

CSc 134

# Database Management and File Organization

## 3. Relational Data Model and Relational Database Constraints

Ying Jin

Computer Science Department

California state University, Sacramento

# Relational Model Concepts

- Relational Model presents a database as a collection of relations.
  - Table :- Relation
  - Row :- Tuple
  - Column header :- attribute

Attribute

## Student

Tuple

Name	SSN	Home Phone
Joe Smith	307-88-2907	602-7765543
Barbara Miller	590-38-6654	422-1076031

# Relational Model

## - Domain

- Domain: A domain D in the relational model is a set of atomic values.
  - Atomic: Each value in the domain is indivisible as far as the relational model is concerned.
- Domain:name, data type, format
- e.g. USA\_Phone\_numbers: A character string of the form (ddd)ddd-dddd, where each d is a numeric (decimal) digit and the first three digits form a valid telephone area code.
- e.g. employee\_age: Possible ages of employee of a company; each must be an integer value between 15 and 80.

# Relational Model

## - Relation Schema

name of the relation

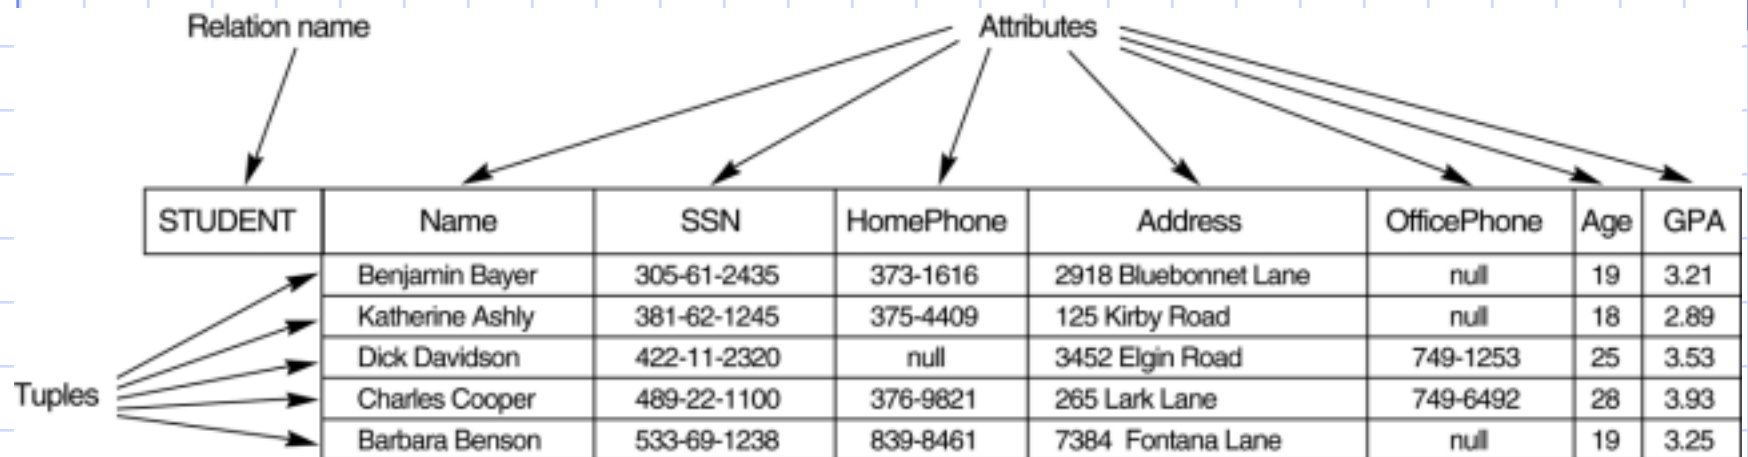
- A relation Schema  $R(A_1, A_2, \dots, A_n)$  is made up of a relation name  $R$  and a list of attributes  $A_1, A_2, \dots, A_n$ 
  - E.g. STUDENT(Name,ssn,phoneNumber)
- Domain of  $A_i$  is denoted by  $\text{dom}(A_i)$   
Degree of a relation: number of attributes  $n$  of its relation schema

# Relational Model

## - relation state

- A relation (or relation state) of the relation schema  $R(A_1, A_2, \dots, A_n)$  is denoted by  $r(R)$ 
  - The relation is a set of of n-tuples  $r = \{t_1, t_2, \dots, t_m\}$ , where each n-tuple  $t$  is an ordered list of values:  $t = \langle v_1, v_2, \dots, v_n \rangle$
  - each value  $v_i$ ,  $1 \leq i \leq n$ , is an element of  $\text{dom}(A_i)$  or is a special **null** value.

unknown or may not apply to a tuple



# Relational Model

## - relation state (Cont.)

- $r(R) \subseteq (\text{dom}(A1) \times \text{dom}(A2) \times \dots \times \text{dom}(A_n))$
- Tuples are unordered in a relation
- A relation cannot have duplicate tuples
- Denote cardinality (number of values) of domain  $D$  by  $|D|$ .
- Maximum number of tuples in  $r(R)$  is  $|\text{dom}(A1)| * |\text{dom}(A2)| * \dots * |\text{dom}(A_n)|$

$r(R)$

The little  $r$ , is the state of the Relation.  $r$  is the captured state at that instance of the relation

$R(\text{color}, \text{gender})$

$\text{color} \mid \text{gender}$

$\text{color} : \{\text{green}, \text{red}, \text{blue}\}$

green, F

$\text{gender} : \{\text{F}, \text{M}\}$

green, M

Total Tuples:

red , F

red , M

$\text{color.size()} * \text{gender.size()}$

blue , F

= 6

blue , M



# Relational Model

## - Attribute value

- Value  $v_i$  in tuple  $t$  for attribute  $A_i$ 
  - $t[A_i]$  or  $t.A_i$
  - E.g. Given tuple  $t = \langle \text{'Joe Smith'}, \text{'307-88-2907'}, \text{'602-7765543'} \rangle$ 
    - $t[\text{Name}] = \langle \text{'Joe Smith'} \rangle$
    - $t.\text{Name} = \langle \text{'Joe Smith'} \rangle$
    - $t[\text{SSN}, \text{Name}] = \langle \text{'307-88-2907'}, \text{'Joe Smith'} \rangle$
    - $t.(\text{SSn}, \text{Name}) = \langle \text{'307-88-2907'}, \text{'Joe Smith'} \rangle$
- An attribute  $A$  of a relation  $R$  can be presented as  $R.A$ 
  - STUDENT.Name

Dot Notation

We've been working on only a single Relation, when working with more than 1, we need 'Constraints'

# Constraints

## - Category

- Constraints on databases can generally be divided into three main categories:
  - **Inherent model-based constraints**
    - constraints that are inherent in the data model
    - e.g.
      - Ordering of tuples in a relation
      - Relational model represents facts about both entities and relationship uniformly a relation
      - A relation cannot have duplicate tuples ( tuh-polls )

Not all relation/entities become tables in the DB

# Constraints

## - Category (Cont.)

- **Schema-based constraints**
  - can be directly expressed in the schemas of the data model, typically by DDL. **Data Definition Language**
- **Application-based constraints**
  - cannot be directly expressed in the schemas of the data model
  - must be expressed and enforced by application program.
- **Another important category of constraints:**  
**Data Dependencies**
  - functional dependencies and multivalued dependencies.

# Schema-based constraints

- Constraints are *conditions* that must hold on *all* valid relation states.
- Domain constraints
- Key constraints
- Constraints on nulls
- Entity integrity constraints
- Referential integrity constraints

# Domain constraints

- Within each tuple, the value of each attribute  $A$  must be an atomic value from the domain  $\text{dom}(A)$ .
- Data type of domain
  - Integer
  - boolean
  - ...

# Key constraints

- SK is a **superkey** of R, if for any two distinct tuples  $t_1$  and  $t_2$  in a relation state  $r$  of R, we have the constraint that  $t_1[SK] \neq t_2[SK]$
- Key constraint, Unique constraint
  - No two distinct tuples in any state  $r$  of R can have the same value for SK.
- e.g. {SSN, Name, Age}

# Key

- A **key** is a minimal superkey – a superkey such that removal of any attribute from K results in a set of attributes that is not a superkey.
- e.g. {ssn}
- A relation schema may have more than one key, each of the keys is called a candidate key.
- e.g. fig

In general, a relation schema may have more than one key. In this case, each of the keys is called a **candidate key**. For example, the CAR relation in Figure 3.4 has two candidate keys: License\_number and Engine\_serial\_number. It is common to designate one of the candidate keys as the **primary key** of the relation. This is the candidate key whose values are used to *identify* tuples in the relation. We use the convention that the attributes that form the primary key of a relation schema are underlined, as shown in Figure 3.4. Notice that when a relation schema has several candidate keys, the choice of one to become the primary key is somewhat arbitrary; however, it is usually better to choose a primary key with a single attribute or a small number of attributes. The other candidate keys are designated as **unique keys**, and are not underlined.

Another constraint on attributes specifies whether NULL values are or are not permitted. For example, if every STUDENT tuple must have a valid, non-NULL value for the Name attribute, then Name of STUDENT is constrained to be NOT NULL.



### Figure 3.4

The CAR relation, with two candidate keys: License\_number and Engine\_serial\_number.

CAR	<u>LicenseNumber</u>	EngineSerialNumber	Make	Model	Year
	Texas ABC-739	A69352	Ford	Mustang	96
	Florida TVP-347	B43696	Oldsmobile	Cutlass	99
	New York MPO-22	X83554	Oldsmobile	Delta	95
	California 432-TFY	C43742	Mercedes	190-D	93
	California RSK-629	Y82935	Toyota	Camry	98
	Texas RSK-629	U028365	Jaguar	XJS	98

In Short both of these candidate keys can be Primary Keys as they are unique

# Primary Key

- Designate one of the candidate keys as the **primary key** of the relation.
- The choice of primary key from candidate keys is arbitrary
- It is better to choose a primary key with *a single attribute* or a *small number* of attributes.
- The primary key attributes are *underlined*.

# Constraints on NULL values

- A constraint specifies that null values are or are not permitted
- e.g. employee Name is constrained to be NOT NULL.

The little button in PHPMYADMIN that says null and is a checkbox.  
Determines whether or not the value can be null, if it is. Throw error.

# Relational Database Schemas

- A **relational database schema**  $S$  is a set of relation schemas  
 $S = \{R_1, R_2, \dots, R_m\}$   
and a set of integrity constraints  $IC$ .
- A **relational database state**  $DB$  of  $S$  is a set of relation states  
 $DB = \{r_1, r_2, \dots, r_m\}$   
such that each  $r_i$  is a state of  $R_i$  and  $r_i$  satisfy the  $IC$ .

# Example of relational database schema

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT\_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS\_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

# One possible database state for the company schema

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
John		Smith	123456789	1966-01-09	751 Fordin, Houston, TX	M	30000	333445555	5	
Franklin		Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888885555	5	
Alice		Zelaya	888887777	1968-01-19	3321 Currie, Spring, TX	F	25000	987654321	4	
Jennifer		Wallace	987654321	1945-05-20	291 Bony, Dallas, TX	F	40000	888885555	4	
Ramesh		Narayan	888884444	1962-09-15	975 Fies Oak, Humble, TX	M	38000	333445555	5	
Joyce		English	453453453	1972-07-30	5831 Res, Houston, TX	F	25000	333445555	5	
Ahmed		Jabbar	987897687	1960-03-29	980 Calos, Houston, TX	M	25000	987654321	4	
Jama		Boj	888885555	1937-11-10	480 Stone, Houston, TX	M	55000	null		

DEPT_LOCATIONS	DNUMBER	DLOCATION
	5	Houston
	4	Stafford
	4	Dallas
	4	Stafford

DEPARTMENT	DNAME	DNUMBER	MGSSN	MGSTARTDATE
Research		5	333445555	1986-06-22
Administration		4	987654321	1995-01-01
Headquarters		1	888885555	1981-05-19

WORKS ON	ESSN	ENO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	888884444	3	40.0
	453453453	1	20.0
	453453453	2	20.0
	333445555	2	10.0
	333445555	3	10.0
	333445555	10	10.0
	333445555	20	10.0
	888887777	30	30.0
	888887777	10	10.0
	987897687	10	35.0
	987897687	30	5.0
	987654321	30	20.0
	987654321	20	15.0
	888885555	20	null

PROJECT	PNAME	PNUMBER	PLOCATION	DNUM
	ProductX	1	Dallas	5
	ProductY	2	Stafford	5
	ProductZ	3	Houston	5
	Computerization	10	Stafford	4
	Reorganization	20	Houston	1
	Newproducts	30	Stafford	4

DEPENDENT	ESSN	DEPENDENT NAME	SEX	BODATE	RELATIONSHIP
	333445555	Alice	F	1986-04-05	DAUGHTER
	333445555	Theodore	M	1983-10-25	SON
	333445555	Joy	F	1958-05-03	SPOUSE
	987654321	Abner	M	1943-02-28	SPOUSE
	123456789	Michael	M	1988-11-04	SON
	123456789	Alice	F	1986-12-30	DAUGHTER
	123456789	Elizabeth	F	1957-05-05	SPOUSE

# Valid /invalid state

- When we refer to a relational database, we implicitly include its *schema* and its *current state*.
- A database state satisfies all the constraints in IC is called a **valid state**.
- A database state does not obey all the integrity constraints is called an **invalid state**.

# Entity integrity constraint

- Entity integrity constraint: No primary key value can be null Duh. If it was null how would we pull it back from the DB
- Because the primary key value is used to identify individual tuples in a relation.
- Involve a *single* relation



# Referential integrity constraints

- Specify a *relationship* among tuples in two relations: the **referencing relation** and the **referenced relation**.
- Informally:
  - refer to an existing tuple

# Foreign Key

- A set of attributes FK in relation schema R1 is a **foreign key** of R1 that **references** relation R2 if it satisfies two rules:
  1. The attributes in FK have the same domain(s) as the primary key attributes PK of R2
  2. A value of FK in a tuple t1 of the current state r1(R1) either occurs as a value of PK for some tuple t2 in the current state r2(R2), or is NULL.

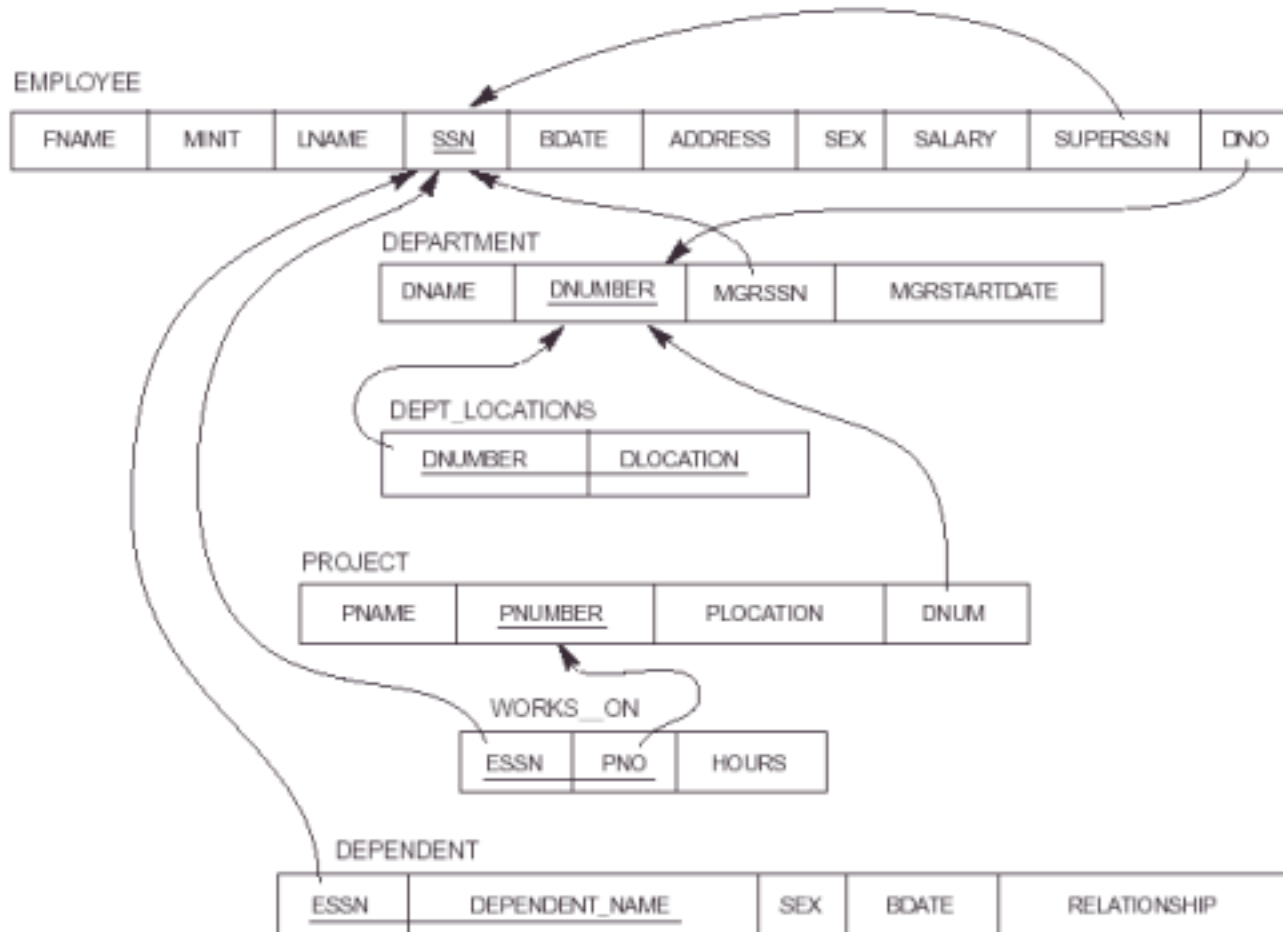
$t1[FK] = t2[PK].$

t1 references or refers to t2.

# Referential integrity constraint definition

- If the two conditions hold, the referencing integrity constraint from  $R_1$  to  $R_2$  is said to hold.
- A referential integrity constraint can be displayed in a relational database schema as a directed arc from  $R_1.FK$  to  $R_2$ .

# Referential integrity constraint example



dno=1?  
dno=null?

# Refer to its own relation

- A foreign key can refer to its own relation.
- e.g. superssn

# Application-based constraints

- Semantic integrity constraints
  - e.g. “The salary of an employee should not exceed the salary of the employee’s supervisor”
- Constraint specification language
  - e.g. trigger, assertions
- Check within application programs

# Update Operations on Relations

- INSERT a tuple.
- DELETE a tuple.
- MODIFY a tuple.
- Integrity constraints should not be violated by the update operations.
- Updates may *propagate* to cause other updates automatically. This may be necessary to maintain integrity constraints.

# Update Operations on Relations (Cont.)

- In case of integrity violation, several actions can be taken:
  - Cancel the operation that causes the violation (REJECT option)
  - Perform the operation but inform the user of the violation
  - Trigger additional updates so the violation is corrected (CASCADE option, SET NULL option)
  - Execute a user-specified error-correction routine



# Constraint violation

## - Insert

- Insert can violate \_\_\_\_\_
  - Domain constraints
  - Key constraints
  - Entity integrity constraints
  - Referential integrity constraints
- Reject the insertion in case of constraint violation

# Constraint violation

## - Delete

- Can violate referential integrity
- In case of violation
  - Reject the deletion
  - Attempt to cascade the deletion
  - Modify the referencing attribute values the cause the violation
    - Set to null
      - foreign key is part of the primary key.
    - Change to reference another valid tuple
  - Specify it in DDL

# Constraint violation

## - Update

- Modify neither a primary key nor a foreign key
  - Check new value in the correct domain
- Update a primary key or a foreign key
  - Delete + Insert
  - Can use DDL to specify how to handle update



These slides are based on the textbook:

R. Elmaseri and S. Navathe, *Fundamentals of Database Systems*, 6th Edition, Addison-Wesley.  
Chapter 7.